

Police are less likely to respond to requests for help from minorities: Field experiment evidence of police discrimination[‡]

For latest draft

Garrett Olson Stanford
Department of Economics, 736 PLC
University of Oregon
Eugene, OR 97403-1285
gos@uoregon.edu

November 2022

Abstract

Numerous high-profile incidents have led to accusations that police departments struggle with equity and accountability. I use an experiment to test both. In the form of a correspondence study, I send emails to more than 2,000 U.S. police departments requesting information about how to lodge a complaint against an officer. Manipulating the names of the email sender, I compare department response rates across race/ethnicity (Black, Hispanic and White) and gender (female and male). I find that departments are less likely to respond to emails signed with Black and Hispanic names. Differences in response rates become more pronounced when I interact gender with race/ethnicity. These differences exacerbate a low overall response rate of 67.4 percent. I find little evidence that department size or the local population are correlated with response rates. Results from this experiment support the accusations that policing suffers from issues of bias and transparency.

Keywords: Policing, Police Discrimination, Police Accountability, Racial Justice, Criminal Justice, Field Experiment

JEL Classification: C93, J15, J16, K40

*I thank Ed Rubin, Jon Davis and Trudy Ann Cameron for invaluable input and guidance. I also thank my peers at the University of Oregon Microeconomics Group for helpful feedback. Finally, I am especially grateful to Joe Mitchell-Nelson for the inspiration and encouragement to undertake this project.

[‡]This study was preregistered with the AEA RCT registry. The pre-analysis plan can be found [here](#).

1 Overview

The Department of Justice reports that there are close to 12,000 local police departments across the United States. For many local governments, these law enforcement agencies (LEAs) are one of the highest-funded agencies.¹ The primary objective of policing is to prevent crime and enforce laws, but the role that LEAs play in American society has expanded over time, both in terms of (1) the civic duties assigned to these agencies, as well as (2) total personnel and infrastructure (Soss and Weaver (2017)). Today, law enforcement officers (LEO) are often expected to take on numerous other social responsibilities.² Given the magnitude of responsibility, public funding, and power invested in LEA, it is imperative to social welfare that LEA perform their duties with an adequate degree of equity and accountability.

Research suggests that police can be effective in reducing crime (e.g., Chalfin et al. (2021), Chalfin and McCrary (2018), Chalfin and McCrary (2017), Cheng and Long (2018), Evans and Owens (2007), Mello (2019), Weisburd (2019b), Weisburd (2021)). However, there is a longstanding debate concerning the impacts that police practices may have on social welfare. Of particular concern is the presence of bias in police behavior—especially racially motivated bias. In 2020, protests over racially biased policing broke out across the nation after George Floyd was killed by officers in the Minneapolis police department. Despite the constant presence of police reform in the national dialogue and despite alarmingly frequent anecdotes reported in the media, few studies exist that causally document bias in the police force (Smith et al. (2017)).

In this paper, I causally estimate the effect of bias on police transparency. To identify this effect, I conducted a field experiment on a sample of 2,134 U.S. police departments. Posing as a fictitious citizen, I emailed each department and requested help making a complaint about an officer in the department. The emails I sent department were identical except for randomly assigned race (Black, Hispanic and White) and gender (male and female) implied by the fictitious sender's name.³ Police departments responded to 67.4 percent of the requests. Response rates for emails signed with Black or Hispanic names were both 10 percentage points (pp) lower than the response rates for emails signed with White names—differences that are both significant at the 1% level. White male emails received the highest rate of response, at 75.8%, which was marginally higher than the response rate for White females. Black and Hispanic males response rates were, respectively, 13.9 and 15 pp lower than White males, and were marginally lower than Black and Hispanic female response rates.

A growing body of research highlights the disproportionate burden that policing can place on people of color. This research covers a variety of contexts: general arrest rates (e.g., Bulman (2019)); predatory fines and asset forfeitures (Makowsky et al. (2019), Sances and You (2017), Shoub et al. (2021), Singla et al. (2020), Su (2020), Su (2021)); police stops and searches (e.g., Antonovics and Knight (2009), Bandes et al. (2019), Feigenberg and Miller (2021), Gelman et al. (2007), Goel et al. (2016), Pierson et al. (2020)); traffic enforcement (e.g., Goncalves and Mello (2021) and West (2018)); and use of force (e.g., Edwards et al. (2019), Fryer (2020), Nix et al.

¹The Census Bureau's financial data reveals, that in 2017, state and local governments spent close to \$200 billion on law enforcement. **cite**

²LEO responsibilities include but not are limited to: responding to mental health crises; serving as first responders to traffic accidents; performing search and rescue; and assisting other social service providers (e.g., Child Protective Services and government sponsored affordable housing agencies). Additionally, LEOs are posted to guard a wide array of public institutions (for example, public parks, schools, hospitals, and airports).

³In this paper I will use "race" as a catchall term for both race (Black and White) and ethnicity (Hispanic). I am sensitive to the distinction of race and ethnicity, and my choice to collapse the distinction in this study is only done for simplicity's sake.

(2017), Ross (2015)). The extent to which racially biased policing can unfairly impact individuals, and sometimes entire communities, can be quite significant. For instance, Edwards et al. (2019) find that police use of force is one the leading causes of death for young men of color in the United States. Most studies that address biased policing, while descriptively informative, are unable to attribute causality. Biased policing centers on the premise that police interact with citizens at a different frequency depending on the sociodemographic characteristics of citizens. Further, conditional on an interaction taking place, policing may still differ across groups. Establishing causality in either of these contexts is difficult. First, differences in frequency of interaction does not necessarily reflect biased policing. There is the possibility of a selection issue. If sociodemographic groups participate in criminal activity at different rates, then unbiased policing will still result in heterogeneous rates of police-citizen interactions across sociodemographic groups (Fridell (2017)). Second, measuring biased policing by comparing outcomes for citizens during an interaction with police does not permit causal inference. If the motivation for police initiating an interaction with a citizen is predicated on bias and outcomes for all police-citizens interactions are similar, then this form of analysis will obscure the presence of biased policing (e.g., Knox et al. (2020) and Ross et al. (2018)). Thus, researchers remain divided on the existence and extent of biased policing (Smith et al. (2017) and Fridell (2017)).⁴ Furthermore, while some researchers employ research designs that permit causal inferences, the vast majority of these studies rely on data supplied by the police themselves. Such reliance on police-reported data can lead to inconclusive or incorrect conclusions if police departments strategically or unintentionally misreport (e.g., Luh (2019)).

The challenge of causally identifying discrimination is not unique to the context of law enforcement. Over the last decade, correspondence studies, a type of randomized controlled trial (RCT), have become an increasingly popular tool for researchers studying the presence of discrimination (Gaddis (2017a), Bertrand and Duflo (2017)).⁵ This paper describes one of the first studies to use a correspondence study to identify the presence of racially and gender-motivated bias in policing, helping address the dearth of evidence concerning causally identified bias in policing.⁶

I make a number of advancements to the extent research on policing with this study. First, by using a correspondence study, I overcome two of the main challenges of studying discrimination in the context of law enforcement: (1) finding causal estimates (as opposed to associations) and (2) avoiding potentially compromised self-reporting of data collected and/or provided by law enforcement agencies. Estimates of properly randomized correspondence studies can be reasonably assumed to be causal. As mentioned above, the primary obstacles to causal inference in the context of biased policing arise from (1) selecting into criminal activities and (2) police discretion of with who they interact. I avoid these challenges by creating a citizen-initiated police interaction that is not predicated on a crime taking place and by designating my outcome of interest as the police departments decision to interact.

Avoiding the use of data provided by police departments has a number of significant advantages. First, departments, either intentionally or inadvertently, can have ongoing difficulties in reporting accurate data (e.g., Luh (2019)). Second, police data can be a

⁴While limited, research has made efforts to address these challenges and find evidence of biased policing (e.g., Antonovics and Knight (2009), Gelman et al. (2007), Ross (2015) and West (2018)).

⁵In a correspondence study, individuals (often fictitious) who are identical for all observable characteristics, other than the characteristic of interest, apply for a job, service or good. The researcher then examines whether the experimentally varied characteristic of interest has an effect on the outcome of the application (Bertrand and Duflo (2017)). In the case of this experiment, email is used instead of the traditional approach of “snail mail,” and—as explained below—a request is made for assistance from police departments, as opposed to applying for jobs or making purchases.

⁶The only other correspondence study examining racial bias in law enforcement is Giulietti et al. (2019)—discussed in more detail below.

product of subjective reporting by individual police officers and department-specific classification standards.⁷ Even when police officers accurately and honestly record officer conduct, decisions made in the heat of the moment during a citizen-officer interaction could influence how events are recorded. Finally, departments may be unwilling to disclose what might be deemed sensitive information.⁸

Second, in addition to estimating a measure of discrimination, this study also contributes to our understanding of police accountability. Response rates of law of enforcement agencies to citizens' requests for any type of assistance is a coarse measure of accountability. By asking departments for help in making a complaint about an officer in their department, I test explicitly test the willingness of departments to hold their officers accountable—an essential concern for policy makers interested in reforming law enforcement in the United States. I am also able to examine how this measure of accountability varies by a citizen's race/ethnicity and gender.

Third, by using a nationally representative sample of police departments, inferences made in this study, by design, are generalizable to policing across the United States. I use a sample of over 2,000 police departments, representing all states except Hawaii.⁹ As a result, the measures of biased policing and accountability from this study are representative of policing in the United States, rather than a specific state, county or city. Consequently, the results reflect systemic behavior, rather than specific department cultures. While both are important for reform, finding evidence of policing issues at the national level make a stronger case for police reform. **[Note: Not exactly sure what I want to say here. But I think it is important to point out that these results are for the whole country, not just “racist Oregon”.**

Finally, this study is that appears to be the first RCT in the context of law enforcement to study discrimination against Hispanic citizens and across genders (male and female). The majority of research concerning biased policing concerns itself with differences between White and Black citizens. However, as [Weitzer \(2014\)](#) points out, considering the growing population of Hispanic Americans, the lack of research on police-citizen interactions for Hispanics is “particularly puzzling.” I also include the dimension of gender in this study. On its own, it is important to understand if police discriminate against a particular gender. Additionally, many of the stereotypes that may motivate biased policing, frequently include the intersection of race/ethnicity and gender. By including race, ethnicity, and gender in my study, I can test if there is any correlation between these problematic stereotypes and policing.

2 Experimental Design and Data

2.1 Experiment

In this section, I describe the design, creation and implementation of the correspondence study. The objective of the study is to test whether police departments exhibit signs of racial/ethnic or gender discrimination. The design of this study, in broad strokes, is to collect contact information for a sample of police departments and then contact the departments using identities that I created. I preregistered this experiment at the AEA RCT Registry, and the pre-analysis plan can be found [here](#).

⁷(For instance, [PolicingProject.org](#) describes the discrepancies in officer-initiated stop report requirements across states. [Reveal-News.org](#) finds that the Washington D.C. police department has a comparatively liberal definition of “resisting arrest”.)

⁸[Weisburst \(2019a\)](#) does not find evidence of racially biased policing in Dallas but hypothesizes that the department's willingness to disclose its data to researchers might stem from the fact that the Dallas police do not appear to have a problem with biased policing.

⁹Hawaii's exclusion was a result of random selection. Hawaii only has four distinct police departments.

Police Department Selection: The police departments included in this study are a stratified random sample. For inclusion in the study, I required a police department to serve a local government (i.e. no state police) and serve a population of at least 7,500 people. To generate my sample, first, I randomly sampled the US Census’s universe of local governments in batches of 1,000.¹⁰ From the sample of local governments, I searched the internet for an email address for the corresponding police department; I conducted a unique search for each department (i.e. I did not use a LEA directory). Some local governments did not have their own police departments, and some police departments did not have publicly available email addresses. In both cases, I recorded the issue and dropped the local government from the study. When a department had multiple email addresses publicly available, priority was given first to the general department email, then to the police chief and then to the next-highest-in-command officer. In the end, I selected 2,134 departments to receive emails, representing 49 states.¹¹ Please refer to appendix A for more details.

Identity Creation: I use the names of the email senders to signal race and gender. I created six broad categories of identities for this study: Black female, Black male, Hispanic female, Hispanic male, White female and White male. Each identity is represented by 60 unique first-last name combinations. I selected last names from the “Frequently Occurring Surnames in the 2010 Census” dataset for this study. I selected last names that were both racially distinctive and commonly occurring. To select racially distinctive names, I found names that are highly concentrated in one racial group—this requires that the name is both common for a particular race and uncommon for other races. However, in some cases, the most racially distinctive names were also names not commonly used in the United States. I avoided using very uncommon surnames to avoid the suspicions of the police departments. I constructed a simple search equation to find names that were sufficiently racially distinctive, while also reasonably common.¹² I selected six last names for each race. I referred to Gaddis (2017b) and Gaddis (2017a) to select the first names. Motivated by the frequent use of names as signals for race in audit studies, Gaddis conducts two experiments that explicitly test which first and last names are racially distinctive. In the experiments, subjects are asked which race they associate with a particular name. This experiment is conducted for names commonly used to represent Black people (Gaddis (2017b)) and Hispanic people (Gaddis (2017a)) in audit studies. I chose the ten most racially distinctive first names for the respective identities from these two studies. In total, I created 360 unique names (6 identities \times 6 last names \times 10 first names).¹³ After selecting the names for each identity, I created a unique email address for each last name used in the study. I then created a unique email address profile for each identity, so that police departments would see the full name of the identity in place of or in addition to the identity’s email address in their inboxes (e.g., *Claire Olson* <olson.2922@mail.com>).¹⁴

Email: Each department received one email from a single randomly assigned identity. All emails were identical with two exceptions (1) the name of the email sender and (2) the sign-off used in the email. I varied the sign-off to test whether the tone of the email had an impact on police behavior. I decided to use the sign-off as the control, because it influences the tone of the email but minimally changes the content of the email. The sign-off was randomly assigned across emails. The name of the email sender is used twice in each email

¹⁰I filtered the universe of local governments to exclude states, counties and all governments with populations less than 7,500.

¹¹Hawaii was not represented in this study. Its exclusion was an unintentional result of the sampling process.

¹²Please refer to Appendix B for details on the names used in this study.

¹³I dropped six of the unique names. These names have celebrity notoriety and during the testing process, respondents noted that they strongly associate these names with the celebrities with the same name (e.g., *Tyra Banks*).

¹⁴Please refer to Appendix C for details concerning the specific email addresses used.

to increase salience. Each email had the following format:

From: *Full name* <lastname19xx@mail.com >

Subject: *Complaint Assistance*

Body:

X Police Department,

My name is *first name* and I am interested in filing a complaint against an officer in your department. I am not sure what to do, and would like to request information on how to make a complaint. Can you please send me this information?

Sign off

Full name

The *italicized* words indicate that these words changed across emails. As seen above, I created profiles for the email accounts so that departments would see an identity's full name twice, and an identity's first name thrice. Police departments were addressed directly—without, for example, a “Hello”—because I found in the testing process that inclusion of a salutation increased the chances of the email being marked as spam. The sign-off varied between a cordial sentiment (“*Thank you!*”) and a curt sentiment (“*Sincerely,*”). Appendix D includes images of example emails, as well as other details on the design of the email template.

Timing: I conducted the study over a ten-week period, from late June 2022 to late August 2022. Each week, I sent roughly 210 emails, split across Monday, Tuesday and Wednesday. I sent all emails at approximately 9 a.m. local time for each police department. I rolled the experiment out over ten weeks to minimize the chance that a single event—in the national or regional news—compromised the generalizability of the results. Splitting the emails across days of the week merely reduced the logistical difficulty of sending the emails. Thursday, Friday and weekend days were not used so that departments had at least two full days to respond before the weekend.

Treatment Assignment The “treatment” here is the identity (race and gender) that each department sees. Treatment was first stratified by week and then by departments’ states, so that the number of departments by state are balanced during each week. Treatment was then randomly assigned across departments within each week-state stratum. Appendix E details on the treatment assignment process.

2.2 Data

[Note: I have not yet included citations for the data sources, but plan on doing so.]

I use several additional datasets in this study. As mentioned, I used data from the US Census to create a pool of local governments in the department-selection process. I limited the local governments eligible for inclusion in the study to exclude state and county governments and governments with populations less than 7,500 residents.¹⁵ I then matched selected departments with police department directories from OpenPolice.org and ICPSR. Department directories added information about agency location and unique identification numbers.¹⁶

¹⁵During the collection process for police department email addresses, I accidentally included 117 departments in communities with populations less than 7,500. These departments were included in the study.

¹⁶I collected the *Originating Agency Identifier* (ORI) numbers for all departments that have ORI. See [Office of Justice’s](#) explanation for details.

The study includes data for several other observable department characteristics. These characteristics, *ex ante*, seemed to be potentially important determinants of the response behaviors of police departments: numbers of officers and civilian employees for each department; county-level income information; and county-level racial/ethnic composition. Employee counts for each department were provided by UCR data codified by Jacob Kaplan (**add citation**). The UCR dataset includes employee counts through 2020. However, a handful of departments are missing data for 2020. Where available, I use the most recent employee count since 2010 (231 departments). If a department does not have an employee count after 2010, I record that department's employee count as missing (29 departments). I take income and race data from the 2019 American Community Survey (ACS). In the 1-year ACS data, 148 counties are missing data for median income of Black households and 55 counties are missing data for median income for Hispanic households. Police departments selected for the study are associated with governments smaller than counties. However, it is unclear with exactly which population each department would interact. If I use data with too precise geography (e.g. the zip codes of the departments), I run the risk of mischaracterizing the local conditions of a department. Accordingly, I use county-level data to characterize the economic and racial composition of a department's local area. I sacrifice precision with this approach, but avoid inaccuracy.

Table 1 shows relevant department characteristics. Column 2 of the table is the mean value for each different characteristic for departments that received emails from White-male identities. Columns 3 through 7 are the differences between the White-male mean value and the other identities. Table 1 confirms that the treatment was successfully randomized across the most obvious department characteristics relevant to this study. Only one of the 85 differences throughout the rows and columns is statistically significant, and only at the 10% level (*Pop. % Black (county-level)*).

Table 1: This table compares mean values of geographic and police departmental characteristics across each identities. Column 1 delineates which variable is being compared. Column 2 displays the mean value of the variables for departments that received emails from White male identities. Columns 3 through 7 show the difference between the value in Column 2 and the mean value for the other five identities. For example, the average county median income for departments that received emails from White male identities is \$66,700. In comparison, the average county median income for departments that received emails from Black female identities is \$100 lower, with a standard error of \$1,300. The absence of statistical significance in the table reflects that the variables in the table are not correlated to treatment assignment.

	Putative Identity					
	White Male (n = 359)	White Female (n = 352)	Hispanic Male (n = 350)	Hispanic Female (n = 361)	Black Male (n = 358)	Black Female (n = 354)
	(Mean)	Differential				
Income (county-level)						
Median Income all HH (hundreds of dollars)	\$667	-0.3 (13)	0.7 (13)	-5 (13)	3.4 (13)	-1 (13)
Median Income Black HH (hundreds of dollars)	\$477	-19 (15)	5 (15)	-1.7 (15)	6.3 (15)	7.6 (15)
Median Income Hispanic HH (hundreds of dollars)	\$531	-6.8 (10)	3.3 (10)	-6.8 (10)	-2.7 (10)	4.8 (10)
Median Income White HH (hundreds of dollars)	\$733	9.6 (14)	6.8 (15)	-12 (14)	3.4 (14)	3.5 (14)
% Pop. in poverty	12	0.39 (0.36)	0.40 (0.36)	0.12 (0.35)	-0.11 (0.35)	0.21 (0.36)
% Black pop. in poverty	22	0.90 (0.76)	-0.07 (0.77)	-0.20 (0.76)	-0.61 (0.76)	0.33 (0.76)
% Hispanic pop. in poverty	19	-0.04 (0.56)	0.25 (0.57)	0.13 (0.56)	0.00 > (0.56)	-0.16 (0.56)
% White pop. in poverty	9	0.06 (0.27)	0.08 (0.27)	-0.06 (0.27)	-0.17 (0.27)	0.02 (0.27)
Population						
Local government pop. (hundreds)	491	41 (128)	-11 (128)	130 (127)	-104 (127)	17 (127)
Pop. % Black (county-level)	10	1.57* (0.84)	0.19 (0.85)	0.64 (0.84)	0.07 (0.84)	0.84 (0.84)
Pop. % Hispanic (county-level)	14	-0.21 (1.11)	0.72 (1.11)	0.37 (1.10)	0.16 (1.11)	0.33 (1.11)
Pop. % White (county-level)	69	-1.99 (1.48)	-1.145 (1.48)	-1.32 (1.47)	-0.23 (1.47)	-1.67 (1.48)
Pop. % rural (county-level)	21	-2.19 (1.51)	-1.29 (1.51)	-0.69 (1.50)	0.61 (1.50)	-0.43 (01.51)
Department size (# of employees)						
Total employees	128	-1 (43)	-13 (43)	48 (43)	-22 (43)	-8 (43)
Total officers	104	-1 (36)	-12 (36)	40 (36)	-22 (36)	-9 (36)
Total civilian employees	24	0 (8)	-1 (8)	8 (8)	0 (8)	1 (8)

Standard-errors in parentheses. Signif. Codes: *: 0.1

Note: 117 departments served local populations < 7,500. Department size data was missing for 29 departments, and 2020 department size data was missing for an additional 231 departments. 148 observations were missing for median income of Black households and 55 observations were missing for median income for Hispanic households in each department's county.

3 Results

3.1 Summary Statistics

I sent the first emails on Monday, June 27th, 2022 and delivered the last emails on Wednesday, August 31, 2022. In total, I attempted to contact 2,134 police departments. Table 2 summarizes the outcomes of these emails. My final analyses (below) exclude the 37 emails that were denied or failed.¹⁷

Table 2: Emails categorized by outcome. Emails were undelivered (failed or denied) because the police department email address was incorrect or police departments blocked the emails.

Email Outcome	Total	Percent of Total
Sent	2,134	100.00%
Response	1,413	66.21%
No Response	682	31.96%
Multiple Response	207	9.70%
Denied	15	0.70%
Failed	24	1.12%

Table 2 indicates a response rate of 66.21%. If I exclude, the 37 emails that were undeliverable (Denied or Failed), the response rate is 67.45%. This response rate is aggregated across all identities and says nothing about biased policing. However, a response rate of 67.45% for the general populace is a striking result in terms of police accountability. The Denied category in column 1 represents emails that were blocked by police departments. The small number of Denied emails does not cause concern for the validity of the argument. However, it is concerning that some police departments have structured their email server to block a seemingly legitimate email request for help. Of course, the request is part of an experiment, but it is easy to imagine a citizen with a real complaint making a similarly formatted request. The Failed emails are

¹⁷During the experiment, I received “undeliverable” emails from the email server I used. These “undeliverable” emails explained why the email I sent could not be delivered. In some cases, the email address I used for a police department was incorrect or no longer existed; these are the “failed” emails. In other cases, the police department email server blocked my email; these are the “denied” emails.




Figure 1

also a cause of concern in terms of police accountability. Because I manually collected the email addresses for each police department from their website, a Failed email reflects a department's neglect in maintaining updated and accessible contact information.

I summarize department response times in Figure 1. The large majority of responses from police departments occur in the first 24 hours that the email is sent, and 97% of the responses were received within 2 days. In contrast to the low response rate, the expediency of response suggests that departments take the request for help seriously. Combining the two results paints a picture that only some departments are willing to assist in making a complaint against an officer, but the willing departments do so actively.

As mentioned, emails were sent out in batches over the course of 10 weeks to reduce the chances of current events influencing police department response behavior. In Appendix G, Figure B2 depicts the response rate for all identities by week and Figure B1 breaks the weekly response rates down by identity. The figures suggest that response behavior did not change considerably, at least during the 10 weeks of the experiment.

3.2 Main results

The primary focus for this study concerns the effect of racial and/or gender biases on transparency in policing. To do this, I estimate variations of the following equation:

$$Response_i = \beta_1 \Pi(Gender_i = Female) + \beta_2 \Pi(Race_i = Black) + \beta_3 \Pi(Race_i = Hispanic) + FE_{i,t} + \varepsilon_i \quad (1)$$

Where i indexes individual police departments, t indexes the week the email is sent, and r indexes the race of the identity. The primary focus for this study concerns the differences in police department response behaviors to White putative identities and Black/Hispanic putative identities. Accordingly, the omitted identity in the analysis is either White or White male. The main outcome, $Response_i$, for this study is a binary indicator for whether or not a police department responded, in any way at all, to an email. For a department to be recorded as having responded, that department must reply within 4 weeks (28 days).¹⁸ FE represent fixed effects for the week an email was sent and a department's state. Because treatment assignment was stratified on week and state, I include fixed effects for both throughout my analyses. Additionally, I cluster standard errors by week and state.

Table 3 reports the most general analysis of differences in response rates across identities, ignoring potential interactions. Column 1 of Table 3 compares department response rates for emails with Hispanic identities (Hispanic emails) and Black identities (Black emails) to the mean response rate for emails with White identities (White emails). The response rate for White emails is 74.82%. Compared to the White email response rate, the response rate for Black emails is -10.4 percentage points (pp) lower [-16.39, -4.36] and the response rate for Hispanic emails is -10.62 pp lower [-15.70, -5.54]. Both estimates are statistically significant at the 1% level.¹⁹ Column 2 compares department response rates for emails with female identities to the mean response rate for emails with male identities (66.02%). The estimates from column 2 show that females, on average, were 2.23 pp more likely [-3.93, 8.40] to receive a response. However, the difference is not statistically significant. The results in column 3 are very similar to columns 1 and 2: Black and Hispanic emails were significantly (10 pp) less likely to receive a response from a police departments. Because I

¹⁸Emails from departments that are automatically generated acknowledgements only that the treatment email was received were not considered to be responses. I discuss alternative definitions of the response variable below.

¹⁹Comparison of the coefficients of the response rate for Black emails and the response rate for Hispanic emails reveals that the estimates are not statistically significant from each other.

assigned the treatments of gender and race independently to departments, race and gender are uncorrelated, and thus the similarity of the results in column 3 is not surprising.

Table 3: General Results

Dependent Variable: Model:	(1)	Response (2)
<i>Variables</i>		
Black	-0.1038*** (0.0307)	
Hispanic	-0.1062*** (0.0259)	
Female		0.0223 (0.0315)
Reference group mean		
White	0.7482	
Male		0.6602
<i>Fixed-effects</i>		
Week	Yes	Yes
State	Yes	Yes
<i>Fit statistics</i>		
Observations	2,094	2,094
R ²	0.05987	0.04939
Within R ²	0.01162	0.00059
<i>Clustered (Week & State) standard-errors in parentheses</i>		
<i>Signif. Codes: ***: 0.01, **: 0.05, *: 0.1</i>		

3.2.1 Interacted and Weighted Results

Motivate interactions: Intersectionality of race and gender are important considerations. Stereotypes play a big role in biased policing. Black male “super predator” “welfare queen” “Karen.” See males as threatening, but not when they are white. Age is also important, but not salient.

To estimate the effect for specific identities, I use the following:

$$Response_i = \beta_1 White_i \times Female_i + \beta_2 Black_i \times Male_i + \beta_3 Black_i \times Female_i + \beta_4 Hispanic_i \times Male_i + \beta_5 Hispanic_i \times Female_i + FE_i + \epsilon_i \quad (2)$$

Where each β indicates the difference in response rate for an identity from the White male response rate. The omitted identity is *White male*. This choice is made (1) so that estimates compare the groups that are commonly discriminated against (non-White and female) to the group commonly given preferential treatment (White male) and (2) for the easiest interpretation of results as the *White male* identity has the highest response rate (75.78) among the six identities. Column 1 of Table 4 compares response rates across the six different identities (3 race categories \times 2 gender categories). Column 2 reports the results of the same estimation equation as column 1, but weights observations by the local populations of the police departments.

The results reported in Table 3 and column 1 of Table 4 assign all departments equal weight. Consequently, the results shown in those tables are interpreted as response rates for an average police department. In other words, if a citizen were to contact a randomly selected police department, the response propensities listed in Tables 3 and 4 are the average response rates for the identities in this study. However, departments that serve larger populations interact with more citizens and are thus likely receive more requests for assistance. By weighting each observation by that department's local population, the response rates shift the interpretation of the key coefficients from the average department's behavior to what the average citizen should expect to encounter. Column 2 of Table 4 weights the observations by the square root of the local population.²⁰

Table 4: Main Results

Dependent Variable:	Response	
Model:	(1)	(2)
<i>Variables</i>		
Continued on next page		

²⁰The square root of population is used as a weight instead of simply the population because of the large distribution of population sizes. For example, Los Angeles has a population of close to 4 million which is over 200 times as large as the median local population (18,000). However, using the common method of logging the populations would reduce the disparity between populations too much. The log of Los Angeles's population is approximately 15, which comparatively similar to the log of the median local population ($\log(18000) \approx 9.8$).

Table 4 – continued from previous page

White × Female	-0.0294 (0.0476)	-0.0625 (0.0542)
Hispanic × Male	-0.1500*** (0.0354)	-0.1490*** (0.0408)
Hispanic × Female	-0.0929* (0.0423)	-0.1419** (0.0512)
Black × Male	-0.1394*** (0.0377)	-0.1715*** (0.0451)
Black × Female	-0.0970** (0.0397)	-0.1017*** (0.0251)
<i>Weights</i>		
Sqrt of local pop.	No	Yes
<i>Fixed-effects</i>		
Week	Yes	Yes
State	Yes	Yes
<i>Fit statistics</i>		
Observations	2,094	2,094
R ²	0.06211	0.07222
Within R ²	0.01397	0.01616
<i>Clustered (Week & State) standard-errors in parentheses</i>		
<i>Signif. Codes: ***: 0.01, **: 0.05, *: 0.1</i>		

Column 1 gives the percentage-point differential in response rates for the five identities compared to the White male identity. Column 2 reveals that response rates for Black males and Hispanic males were significantly lower than White males and are the lowest among all identities

(-13.94 [-21.33, -6.56]) and Hispanic males (-15.00 [-21.94, -8.06]) are statistically significant at the 1% level and are the lowest among all the identities. In comparison, the response rates is -9.70 [-17.48, -1.92]) for Black females and is -9.29 [-17.57, -1.00] for Hispanic females. The estimates, respectively, are statistically significant at the 5% and 10% level. Testing for equality between the coefficients within each race/ethnicity group between genders finds that the response rates for Black males and Black females are not statistically significant (p value = 0.3119), while the response rates for Hispanic males and Hispanic




Figure 2: Response rate differentials from the response rate for White male identity from the other identities. The mean response rate for White male identities is 75%. Purple indicates unweighted estimates and gold indicates weighted estimates. Weights are determined by local population size for a department.

females are statistically significant (p value = 0.0035). The response rate for White females is -2.94 [-12.27, 6.40] percentage points lower, but is not statistically significant at the 10% level. White females are the only female identity to have a lower response rate than the male complement within each race/ethnicity grouping.

Weighting by local populations increases the disparity in the response rate for White males and all the other identities, except for Hispanic males. The response rate for White females in column 2 (-6.25 [-16.88, 4.37]) is more than double the estimate from column 1, but remains statistically insignificant at the 10% level. Response rates for Black females (-10.17 [-15.08, -5.25]) and Hispanic females (-14.19 [-24.21, -4.16]) both increase in magnitude and statistical significance. The response rate for Black males (-17.15 [-26.00, -8.31]) becomes the identity with the lowest response rate. The response rate for Hispanic males (-14.90 [-22.89, -6.90]) decreases marginally in magnitude, but remains statistically significant at the 1% level. In contrast to column 1 of Table 4, testing for equality between the coefficients within each race/ethnicity group between genders finds that the response rates for Black males and Black females *are* statistically significant (p value = 0.0861), while the response rates for Hispanic males and Hispanic females *are not* statistically significant (p value = 0.8545).

3.2.2 Sample split by Department Size

A potentially important factor in response behavior for a department is the number of employees that the department has at its disposal. Table 5 displays the results of models with the same specifications as Table 4 when the data is split into “small departments” and “big departments”. The size of the department is determined by the median number of total employees for the departments included in the present study. Response rates are lower for White females, Black males and Black females in big departments compared to small departments. The White female response rates for small departments are 1.62 [-11.84; 15.08] (unweighted) or 2.10 [-12.72; 16.93] (weighted) percentage points *higher* than White male response rates. However, the White female response rates for big departments are -7.03 [-21.71; 7.65] (unweighted) or -9.81 [-24.93; 5.30] (weighted) percentage points *lower* than White male response rates. The unweighted differential response rate for Black males is more than twice as small for big departments (-18.92 [-36.25; -1.59]) as small departments (-8.92 [-16.64; -1.20]). When the observations are weighted by population, the differential response rate for Black males from big departments (-21.87 [-38.09; -5.65]) is almost three times as small as

the response rate from small departments (-7.89 [-14.49; -1.28]). In contrast, the response rate for Hispanic males increases with department size, increasing from -17.55 [-26.63; -8.48] to -13.08 [-25.31; -0.84] when unweighted by population size, and increasing from -17.86 [-28.21; -7.51] to -14.43 [-26.25; -2.60] when weighting by population size. The most dynamic identity response rate is for Hispanic females. The unweighted results show that response rate *increases* with departments size from -8.80. [-19.40; 1.80] to -7.30 [-21.00; 6.39] percentage points, but when weighting by population size the response rate decreases with *decreases* with department size from -8.51 [-20.03; 3.01] to -14.79 [-29.95; 0.36] percentage points. Column 3 and Column 6 show the results of explicitly testing for differences in identity response rate estimates across sample sizes. Three of the differences are significant at the 10% level: Hispanic male response rates for the weighted and unweighted results, and the Black female response rate for the weighted results. In the other cases, the change in point estimates were matched with larger standard errors.

Table 5: Sample Split by Department Size

Dependent Variable:			Response			
	Unweighted			Weighted by Population		
Model:	Small Dept. (1)	Big Dept. (2)	P value	Small Dept. (3)	Big Dept. (4)	P value
<i>Variables</i>						
White × Male mean	0.747	0.765				
White × Female	0.0162 (0.0595)	-0.0703 (0.0649)	0.5605	0.0210 (0.0655)	-0.0981 (0.0668)	0.5523
Hispanic × Male	-0.1755*** (0.0401)	-0.1308** (0.0541)	> 0.0000	-0.1786*** (0.0458)	-0.1443** (0.0523)	>0.0000
Hispanic × Female	-0.0892* (0.0469)	-0.0730 (0.0605)	0.1183	-0.0851 (0.0509)	-0.1479* (0.0670)	0.3045
Black × Male	-0.0892** (0.0341)	-0.1892** (0.0766)	0.6592	-0.0789** (0.0292)	-0.2187** (0.0717)	0.8172
Black × Female	-0.0882 (0.0496)	-0.1039 (0.0569)	0.1611	-0.0813 (0.0516)	-0.1132*** (0.0337)	0.06056
<i>Weights</i>						
Sqrt of local pop.	No	No		Yes	Yes	
<i>Fixed-effects</i>						
Week	Yes	Yes		Yes	Yes	

Continued on next page

Continued on next page

Table 5 – continued from previous page

State	Yes	Yes	Yes	Yes
<i>Fit statistics</i>				
Observations	1,052	1,014	1,052	1,014
R ²	0.08835	0.09749	0.09078	0.11283
Within R ²	0.01864	0.01648	0.02000	0.02028
<i>Clustered (Week & State) standard-errors in parentheses</i>				
<i>Signif. Codes: ***: 0.01, **: 0.05, *: 0.1</i>				

4 Discussion

4.1 Results

4.1.1 General Results and Interacted Results Discussed

The results of this field experiment present strong evidence of racially biased policing. When aggregated across genders, compared to White emails, the response rate for Hispanic and Black emails are both 10 percentage points lower and statistically significant at the 1% level (Table 3). The question of gender biased policing is slightly more nuanced. Comparing response rates for all males to response rates for all females, it appears that the police are slightly more likely to respond to requests from females, though the difference is not statistically significant. However, when response rates for each race/ethnicity-gender identity are compared, a different story emerges (Table 4). Perhaps unsurprisingly, White male identities receive the highest rate of responses—a reversal of the female identities receiving preferential treatment result. It is the low response rates for Hispanic males and Black males that drive the lower response rates for all males. Even though, in aggregate, there is not a statistically significant difference in male and female response rates, Hispanic females and Black females response rates are 9 percentage points lower than White males and significant at the 10% and 5% level, respectively. Comparing the results of Table 3 to Table 4 reveals that the *intersection* of race/ethnicity and gender is an important part of the story. Hispanic males receiving the lowest response rates of all groups indicates the importance of expanding police-citizen relationship research to include Hispanic demographics (Weitzer (2014)).

Identifying the mechanism(s) for the hierarchy of response rates of the six identities is beyond the scope of this paper. However, it is worth considering why Hispanic *males* and Black *males* received the lowest

rate of responses, despite White *males* receiving the highest rate of responses. The discrepancy could be explained by the historical narrative of black and brown males being viewed as criminals (e.g., the racist stereotype of the “superpredator”). A common rebuttal to this hypothesis is that these groups might be more likely to participate in crime—echoing the challenge researchers run into of separating biased policing from different levels of participation in criminal activities amongst different ethnic/racial groups. However, in the context of the present study, a crime is not being committed. Consequently, whether it is due to racist stereotypes, historic crime patterns, or a different mechanism, Black and Hispanic males are not treated the same way as their White counterparts.

An alternative explanation to the discrepancy across the six identities could be that departments hypothesize that the nature of the complaint might be different or that the likelihood of a formal complaint being pursued differs across groups. For instance, research suggest that police are more likely to use force with non-White citizens (e.g, [Edwards et al. \(2019\)](#), [Fryer \(2020\)](#), [Nix et al. \(2017\)](#), [Ross \(2015\)](#)). The lower response rates from the present study might reflect that departments think that complaints coming from Hispanic and Black males are more likely to concern excessive use of force from one of their officers, and that type of complaint is more damaging to the department than other types of complaints.

In studies that examine racial discrimination, the question of statistical discrimination versus racial bias often comes up. In the case of this study, statistical discrimination on the grounds of different levels of criminal activity has already been addressed. However, another common refrain is that the discriminating against Black or Hispanic populations might be a symptom of these populations more frequently having lower socioeconomic statuses (SES) and police departments are discriminating on SES rather than race. For instance, [Giulietti et al. \(2019\)](#) take measures to separate the two in their correspondence study. However, if researchers are unable to disentangle the effect of minority status and lower SES on discrimination then in practice why does it matter? Impact over intent. [Bertrand and Duflo \(2017\)](#) argue that it is possible that police who disproportionately target non-White groups are engaging to some degree in both statistical targeting and biased policing ([Bertrand and Duflo \(2017\)](#)). [Tilcsik \(2021\)](#) argues that the idea of statistical discrimination “can lead people to view social stereotyping as useful and acceptable and thus help rationalize and justify discriminatory decisions.”

4.1.2 Weighted and Department Size Results Discussed

Police department response rates might also be driven by other factors. Size of a department (number of employees) and the population that a department serves could potentially affect response rates from departments. For instance, bigger departments might be able to staff personnel solely responsible for replying to requests for help making complaints. Conversely, smaller departments might be more sensitive to officer-complaints if they are familiar with all officers in the department. Departments that serve larger populations could mean that departments have more requests to fulfill. On the other hand, if departments serve smaller populations they might be more suspicious of the genuineness of the email they receive.²¹

FIX Comparing Columns 1 and 2 in Table 4 it is apparent that weighting by the population of the police department's jurisdiction exacerbates the differences in response rates (with the exception of Hispanic male response rates). One possible explanation for these results is that departments serving larger populations are more likely to discriminate. To test this interpretation, departments are separated into five bins determined by the quintiles of the local populations of the departments included in the study. Model 1 from 3 is re-estimated, but this time interacting both *Hispanic* and *Black* with the five population bins. Figure 3 shows the results of this exercise. Figure 3 reveals no clear pattern in the relationship between population size and response rate. Black identity response rates are lowest for the largest and smallest quintiles, but the pattern does not hold for Hispanic identity response rates.

The relationship between police department local population and response rate is ambiguous. Notwithstanding a clear relationship between local population and propensity to discriminate, heterogeneous response rates across population sizes suggest that studies that restrict their area of focus to a limited number of local governments, or at least similarly sized populations, may not be able to extend their results to smaller (or larger) populations. Furthermore, the higher differentials in response rates between White and non-White response rates when weighting by population (Table 4) indicate, in a crude measure, that more people are discriminated against than is evident by finding the level of discrimination for the average police department.²²

²¹In the response from more than one department the police officer indicated that they checked the police logs and had no record of an interaction between one of their officers and a person that matched the name in the email they received.

²²The design of this study only sends one email to each department, so a department responding to a White email does not necessarily mean the same department would *not* respond to a non-White email. Conversely, a department not responding to a non-White email does not necessarily mean the same department *would* respond to a White email. However, higher levels of discrimination for departments




Figure 3: Response rates differentials for Black and Hispanic identities from White identities separated into bins determined by local population size.

A variable strongly correlated with local population size is the number of employees working for a police department. Mechanically, as local populations increase, so do the sizes of departments. For the departments include in the present study there are on average 3 employees for every 1,000 residents. Results from Table 5 suggest that bigger departments are more likely to discriminate against Black identities and White female identities, and discriminate less against Hispanic identities. When the results are weighted by population in Columns 3 and 4 in 5, the discrepancies grow, which makes sense given that agency size is correlated with population size.²³ To disentangle the effect of agency size and population size on response rate given their the two sizes correlations, the same exercise from Figure 3 is done. This time the quintiles are determined by the number of employees divided by local population size—a per capita measurement. Figure 4 displays the results of this exercise.

No clear pattern emerges for Hispanic or Black identity response rates. Compared to Figure 3, Figure serving larger populations do suggest that a larger share of the population might be discriminated against than is indicated by the unweighted results.

²³One curious result of Table 5 is comparing the point estimate for Hispanic females from Column 2 and Column 4.

Response Rate differential from White grouped by dept. employee per 1,000 residents

Group population size: 1 = [0 : 1.66], 2 = [1.66 : 2.03], 3 = [2.03 : 2.41], 4 = [2.42 : 2.95], 5 = [2.95 : 15.7]




Figure 4: Response rates differentials for Black and Hispanic identities from White identities separated into bins determined by local population size.

4 has smaller differences for the first bin. However, in both Figures the biggest bin, bin 5, exhibits large differences for Hispanic and Black response rates. It is conceivable that the low response rates for departments with big populations might be attributed to departments being overextended and thus less capable of responding to requests for assistance. However, the highest employee per capita departments exhibiting the highest degree of discrimination runs counter to that argument. It is concerning that departments most capable, in terms of employees per capita, to respond to requests are most likely to discriminate.

4.2 Accountability

The primary question of this study seeks to answer whether police departments discriminate on race, ethnicity or gender. The design of the study also emphasizes another important topic for policy makers interested in reforming police practices: accountability. By employing a correspondence study where citizens ask for assistance making a complaint against an officer, police departments are being asked to make a decision to respond to an inquiry based on a citizen's race/ethnicity and gender. However, they are also being assessed

for their willingness to assist a citizen who is attempting to hold one of their officers accountable. In existing literature, the only prior correspondence study that includes law enforcement agencies is [Giulietti et al. \(2019\)](#). Their correspondence study interacts with a wide range of public institutions, where the list includes sheriffs' offices. In their study, the authors email the various public institutions with benign requests for relevant information, varying the identity of the citizen asking for information (where they use two distinctively black male names and two distinctively white male names). The authors find that these public institutions (ranging from public libraries to county clerks, in addition to sheriffs' offices) are less likely to respond to emails from individuals with distinctively black names. [Giulietti et al. \(2019\)](#) find that response rates for their sheriff's offices are above 53% for White male emails and about 46% for Black male emails. Overall, these response rates are noticeably lower than average response rates in the present study. However, the difference in response rates by race are significantly smaller. One explanation for the difference is that [Giulietti et al. \(2019\)](#) target sheriff's offices instead of local police departments, and sheriffs' offices may face different expectations for accountability. However, an alternative explanation is that a simple request for general assistance is not treated with the same urgency as a request for help in making a complaint against a police officer. When a request for assistance concerns making a complaint, police departments seem to be more responsive, but they also more evidence of racial/ethnic discrimination.

It should be noted that in both studies, [Giulietti et al. \(2019\)](#) and the present study, average response rates are quite low. The average response rate of 67.4% for the present study (with a low of 60.6% for Hispanic males) is concerning. Even the most responded to identity, White males, have a response rate of only 75%. By design, the complaints mentioned in the present study are fictitious. However, in reality, an officer behaving in such a way to warrant a citizen making an effort to file a formal complaint suggests potentially serious misconduct on the part of the officer. If only six out of ten citizens are able even to get assistance making a complaint, much less to have any action taken as a result of their complaint, it is difficult to consider citizen-initiated complaints about police officers as a reliable strategy for holding police officers accountable. This concern is amplified when groups of people who more often interact with police (i.e., non-Whites) are also less likely to be assisted in making a complaint.

4.3 Caveats

This study seeks to understand whether police departments tend to discriminate on the basis of race, ethnicity or gender. The results suggest that the average police department does. There are a few caveats to this study to keep in mind. First, police departments were selected randomly (see Appendix ??). However, for a department to be eligible to be included in the study, it needed to have a publicly available email address. There are likely to be non-random department characteristics that distinguish between departments that make their email addresses available to the public and those that do not. Consequently, this study's results reflect average department behavior only for a certain type of department. Intuitively, it is likely that departments willing to share a contact email might also be more willing to engage with the public. Thirty-nine departments included in the study had contact emails found somewhere other than on the police department website (e.g., the police chief's contact information is posted on the city's website, but not the police department's own website or the department's specific page on the city website). The response rate for emails found this way was almost 20 percentage points lower than the overall mean (47% versus 66%). It is difficult to make clear references from such a small sample, but this difference in response rates suggests that departments with easier-to-find email addresses may be more willing to engage with the public. Second, this analysis does not seek to identify a fixed effect for each police department. While the results clearly demonstrate that police departments in the United States have a higher propensity on average to respond to White emails as opposed to Black or Hispanic emails, the data provide only one observation per department. Thus it is not possible to infer systemic bias within individual police departments. Revisiting this type of RCT with a specific aim of learning more about within-department behavior may be of value in future studies.²⁴ Finally, one must keep in mind that this conclusion pertains to a specific context. This study demonstrates that police departments discriminate on race, ethnicity or gender when contacted via email for help making a complaint against an officer. In other contexts, for example, a police officer's decision to pull over a vehicle, the results of this study may or may not apply.

²⁴Multiple requests to one department may have raised department suspicions about the inquiries.

4.4 Conclusion

This study uses a correspondence study to establish strong evidence for the existence of biased policing in the United States. From a sample of over thousand police departments, departments were 10 percentage points more likely to respond to emails from White identities than to emails from Black or Hispanic identities. When the racial/ethnic identities were interacted with gender, White male identities had the highest response rates and Black male and Hispanic male identities had the lowest response rates—respectively, 13.94 and 15 percentage points lower than the White male response rates. Further analysis of the responses that were received and a follow-up RCT that seeks to examine within-department behavior would greatly complement the striking findings of this study. The low overall response rates and large heterogeneity in response rates across identities are concerning. Low response rates suggest that police departments are resistant to accountability. Furthermore, response rates being significantly lower for emails from non-White identities, suggest that departments are especially unwilling to engage with communities of color. Although the literature remains inconclusive about the existence of biased policing, the results of this study suggest that biased policing does exist, and it is imperative to identify biased policing in other contexts.

References

- Antonovics, K. and Knight, B. G. (2009). A new Look at racial profiling: Evidence from the Boston Police Department. *Review of Economics and Statistics*, 91(1):163–177.
- Bandes, S. A., Pryor, M., Kerrison, E. M., and Goff, P. A. (2019). The mismeasure of Terry stops: Assessing the psychological and emotional harms of stop and frisk to individuals and communities. *Behavioral Sciences and the Law*, 37(2):176–194.
- Bertrand, M. and Duflo, E. (2017). Field Experiments on Discrimination a Laura Stilwell and Jan Zilinsky provided excellent research assistance. We thank Abhijit Banerjee for comments. We are particularly grateful to Betsy Levy Paluck, our discussant, for her detailed and thoughtful review. pages 309–393.
- Bulman, G. (2019). LAW ENFORCEMENT LEADERS AND THE RACIAL COMPOSITION OF ARRESTS. *Economic Inquiry*, 57(4):1842–1858.
- Chalfin, A., Hansen, B., Weisburst, E., and Williams, M. (2021). Police Force Size and Civilian Race. *SSRN Electronic Journal*.
- Chalfin, A. and McCrary, J. (2017). Criminal deterrence: A review of the literature.
- Chalfin, A. and McCrary, J. (2018). Are U.S. cities underpoliced? Theory and evidence. *Review of Economics and Statistics*, 100(1):167–186.
- Cheng, C. and Long, W. (2018). Improving police services: Evidence from the French quarter task force. *Journal of Public Economics*, 164:1–18.
- Edwards, F., Lee, H., and Esposito, M. (2019). Risk of being killed by police use of force in the United States by age, race–ethnicity, and sex. *Proceedings of the National Academy of Sciences of the United States of America*, 116(34):16793–16798.
- Evans, W. N. and Owens, E. G. (2007). COPS and crime. *Journal of Public Economics*, 91(1-2):181–201.
- Feigenberg, B. and Miller, C. (2021). Would Eliminating Racial Disparities in Motor Vehicle Searches have Efficiency Costs? *The Quarterly Journal of Economics*, 137(1):49–113.

- Fridell, L. A. (2017). Explaining the Disparity in Results Across Studies Assessing Racial Disparity in Police Use of Force: A Research Note. *American Journal of Criminal Justice*, 42(3):502–513.
- Fryer, R. G. (2020). An empirical analysis of racial differences in police use of force: a response. *Journal of Political Economy*, 128(10):4003–4008.
- Gaddis, S. M. (2017a). how black are Lakisha and Jamal? Racial perceptions from names used in correspondence audit studies. *Sociological Science*, 4:469–489.
- Gaddis, S. M. (2017b). Racial/Ethnic Perceptions from Hispanic Names: Selecting Names to Test for Discrimination. *Socius: Sociological Research for a Dynamic World*, 3:237802311773719.
- Gelman, A., Fagan, J., and Kiss, A. (2007). An analysis of the New York City police department's "stop-and-frisk" policy in the context of claims of racial bias. *Journal of the American Statistical Association*, 102(479):813–823.
- Giulietti, C., Tonin, M., and Vlassopoulos, M. (2019). Racial discrimination in local public services: A field experiment in the United States. *Journal of the European Economic Association*, 17(1):165–204.
- Goel, S., Rao, J. M., and Shroff, R. (2016). Precinct or prejudice? Understanding racial disparities in New York city's stop-and-frisk policy. *Annals of Applied Statistics*, 10(1):365–394.
- Goncalves, F. and Mello, S. (2021). A few bad apples? racial bias in policing. *American Economic Review*, 111(5):1406–1441.
- Knox, D., Lowe, W., and Mummolo, J. (2020). Administrative Records Mask Racially Biased Policing. *American Political Science Review*, 114(3):619–637.
- Luh, E. (2019). Not So Black and White: Uncovering Racial Bias from Systematically Masked Police Reports. *SSRN Electronic Journal*.
- Makowsky, M. D., Stratmann, T., and Tabarrok, A. (2019). To serve and collect: The fiscal and racial determinants of law enforcement. *Journal of Legal Studies*, 48(1):189–216.
- Mello, S. (2019). More COPS, less crime. *Journal of Public Economics*, 172:174–200.

- Nix, J., Campbell, B. A., Byers, E. H., and Alpert, G. P. (2017). A Bird's Eye View of Civilians Killed by Police in 2015: Further Evidence of Implicit Bias. *Criminology and Public Policy*, 16(1):309–340.
- Pierson, E., Simoiu, C., Overgoor, J., Corbett-Davies, S., Jenson, D., Shoemaker, A., Ramachandran, V., Barghouty, P., Phillips, C., Shroff, R., and Goel, S. (2020). A large-scale analysis of racial disparities in police stops across the United States. *Nature Human Behaviour*, 4(7):736–745.
- Ross, C. T. (2015). A multi-level Bayesian analysis of racial Bias in police shootings at the county-level in the United States, 2011-2014. *PLoS ONE*, 10(11):e0141854.
- Ross, C. T., Winterhalder, B., and McElreath, R. (2018). Resolution of apparent paradoxes in the race-specific frequency of use-of-force by police. *Palgrave Communications*, 4(1):1–9.
- Sances, M. W. and You, H. Y. (2017). Who pays for government? descriptive representationb and exploitative revenue sources. *Journal of Politics*, 79(3):1090–1094.
- Shoub, K., Christiani, L., Baumgartner, F. R., Epp, D. A., and Roach, K. (2021). Fines, Fees, Forfeitures, and Disparities: A Link Between Municipal Reliance on Fines and Racial Disparities in Policing. *Policy Studies Journal*, 49(3):835–859.
- Singla, A., Kirschner, C., and Stone, S. B. (2020). Race, Representation, and Revenue: Reliance on Fines and Forfeitures in City Governments. *Urban Affairs Review*, 56(4):1132–1167.
- Smith, M. R., Rojek, J. J., Petrocelli, M., and Withrow, B. (2017). Measuring disparities in police activities: a state of the art review.
- Soss, J. and Weaver, V. (2017). Police Are Our Government: Politics, Political Science, and the Policing of Race-Class Subjugated Communities.
- Su, M. (2020). Taxation by Citation? Exploring Local Governments' Revenue Motive for Traffic Fines. *Public Administration Review*, 80(1):36–45.
- Su, M. (2021). Discretion in Traffic Stops: The Influence of Budget Cuts on Traffic Citations. *Public Administration Review*, 81(3):446–458.

- Weisburd, S. (2021). Police presence, rapid response rates, and crime prevention. *Review of Economics and Statistics*, 103(2):280–293.
- Weisburd, E. K. (2019a). Police Use of Force as an Extension of Arrests: Examining Disparities across Civilian and Officer Race. *AEA Papers and Proceedings*, 109:152–156.
- Weisburd, E. K. (2019b). Safety in police numbers: Evidence of police effectiveness from federal cops grant applications. *American Law and Economics Review*, 21(1):81–109.
- Weitzer, R. (2014). The puzzling neglect of Hispanic Americans in research on police–citizen relations. *Ethnic and Racial Studies*, 37(11):1995–2013.
- West, J. (2018). Racial Bias in Police Investigations. (October):1–36.

A Online Appendix: Police Department Selection

The selection process for police departments to be included in the study is as follows:

1. From the universe of governments provided by the U.S. Census, I create a list of possible jurisdictions that may have their own police departments. This list excludes state governments, counties, special districts and places with populations less than 7,500.
2. From the “possible department” list, I randomly draw 1,000 jurisdiction names.²⁵
3. Email addresses are then collected from police department websites in these jurisdictions.²⁶
 - Governments without local police departments are dropped.
 - Police departments without publicly available email addresses are documented and dropped.
 - In the cases where there are multiple email addresses the prioritization is given first to (1) the email address for the department in general, and then to (2) the email address specifically for the police chief, (3) and finally to any possible contact (e.g., a community-affairs officer). I document the type of email address ultimately recorded in my database.
4. I repeat Steps 2 and 3 until 2,000 email addresses have been collected.²⁷

Randomization of the department selection process increases the external validity of the study. Requiring that populations served by these police departments are greater than 7,500 increases the plausibility of the existence of the purported email sender as a resident of their jurisdiction.

Include a visual of where the addresses came from

Summary stats for dropped etc departments

A.1 Type of email collected

During the police department email address collection, the “type” of email address publicly available for collection varied from department to department. In this case, “type” refers to who is associated with the

²⁵The target number of departments is 2,000, but to streamline the process, I select possible department cities in 3 batches of 1,000. For each batch of jurisdictions, roughly 60% have a viable police department, police chief or alternate email addresses. For a full explanation of this process see ??

²⁶For an example of how this process works please see ??

²⁷Given that jurisdictions are selected in batches of 1,000, the final number of police department emails collected is 2,135.

email address. For example, for Department **X** the only publicly available email address is for the chief of that department and for Department **Y** the only publicly available email address is for the shift-commander. In this example, I collect the email address for each department, and record that the *email address type* for Department **X** is *chief* and the *email address type* for Department **Y** is *shift-commander*. During the actual department email address collection, frequently departments had multiple email addresses publicly available.²⁸ In the case of the existence of multiple publicly available email addresses, I used a consistent priority list to decide which email address to collect. Prioritization is as follows:

1. Top priority is given to a general department email address. This is done to get the most accurate representation of a department's general behavior.
2. In the absence of a general department email address, priority is given to the chief of police.
3. In the absence of a general department email address or a chief email address, priority is given to the next highest in command officer.
4. In the absence of (1) (2) and (3), the email address for the records department is collected.
5. If none of the above email addresses are publicly available, any email address available on the the police department website is collected.
6. If there are no email addresses available on the department website, a cursory search is performed to find email contacts on other related websites (e.g., the website of the city that a department is located in or the official Facebook page for a department)

Table **B1** displays a few basic statistics concerning the email type for the police departments in the study.²⁹ Column 1 categorizes the types of email addresses collected. Some of the email categories warrant explanation. *Commanding Officer* refers to any email address associated with an individual police officer that is in some managerial position, but isn't the chief of police. *Found* refers to an email address collected from a website or web page separate from a police department's website/web page. This could mean collecting a chief's email address from a city's website directory, even though the chief's email address is not listed on

²⁸In other instances, there were no publicly available email addresses associated with a police department of interest.

²⁹Table **B1** excludes the 39 email addresses that were unable to receive emails. Consequently, Table **B1** only has 2095 observations—2,134 total emails sent less the 39 emails that were undeliverable.

the page designated for police on that city’s website. Another example, of a *Found* email would be a contact email being listed on a department’s official Facebook page, but not on the police department’s website or the police web page of a city’s website. *Accountability* refers to email addresses collected in step (5) of the email-collection-priorities list that are associated with community outreach, community engagement or internal affairs. For instance, a handful of departments’ only publicly available email address was for an officer designated to assisting with complaints against officers in their department.

Table B1: Different “types” of email addresses were collected, out of necessity, for different police departments. Column 1, Email Category, describes the type of email address collected (definitions above). The majority of the email addresses used in this study were *Department* or *Chief* email addresses. Although there is variation in Column 4, the number of observations are quite low.

Email Category	Count	Percent of Total (2095	Response Rate
Department	853	40.72%	0.659
Chief	1063	50.74 %	0.687
Commanding Officer	84	4.01%	0.726
Found	40	1.91%	0.45
Records	37	1.77%	0.865
Accountability	18	0.85%	0.556

As seen in Table B1, the vast majority of email types are *Department* and *Chief*. There is variation in response rate across the email types, but given the low count number for the email types other than *Department* and *Chief*, caution should be used when making inferences from these response rates.³⁰ For *Department* and *Chief* email types, response rates compare similarly—at least when aggregating all putative identities. What can be learned about response rate differences for individual putative identities across the different email types? As described in the main body of the paper, response rates were considerably different across putative identities. Because treatment assignment did not directly account for differing email type, the first step in analysing the putative identity differential response rates across email types is to examine

³⁰Granted the sample size of *Accountability*-type emails is quite small (18), but the low response rate is a bit unexpected, given the nature of the email type.

if there is any correlation between putative identity and email type. Table B2 gives the number count for emails sent with each putative identity for each email type. Table B2 shows that there is some heterogeneity across putative identities and emails types. Table B3 gives a more rigorous analysis of the relationship between the putative identity assigned to a department and the type of email collected from that department.

Table B2: This table shows the count of emails by email address type for each putative identity. It complements B3.

Dependent Variables: Model:	Department (1)	Chief (2)	Comm. Officer (3)	Found (4)	Records (5)	Accountability (6)
<i>Putative Identities</i>						
White × Male	0.4188*** (0.0262)	0.4701*** (0.0267)	0.0484*** (0.0105)	0.0285*** (0.0073)	0.0285*** (0.0070)	0.0057 (0.0049)
<u>Differential from White Male</u>						
White × Female	-0.0304 (0.0373)	0.0488 (0.0379)	-0.0050 (0.0149)	0.0063 (0.0104)	-0.0198** (0.0100)	<0.0000 (0.0070)
Hispanic × Male	0.0156 (0.0373)	0.0285 (0.0379)	-0.0164 (0.0149)	-0.0110 (0.0104)	-0.0168* (0.0100)	0.0001 (0.0070)
Hispanic × Female	-0.0323 (0.0369)	0.0509 (0.0376)	-0.0008 (0.0148)	-0.0145 (0.0103)	-0.0089 (0.0099)	0.0055 (0.0069)
Black × Male	-0.0342 (0.0371)	0.0855** (0.0377)	-0.0199 (0.0148)	-0.0171* (0.0103)	-0.0171* (0.0099)	0.0028 (0.0070)
Black × Female	0.0122 (0.0372)	0.0098 (0.0378)	-0.0082 (0.0149)	-0.0199* (0.0103)	-0.0026 (0.0100)	0.0087 (0.0070)
<i>Fit statistics</i>						
Observations	2,095	2,095	2,095	2,095	2,095	2,095
R ²	0.00188	0.00325	0.00146	0.00471	0.00330	0.00125
Adjusted R ²	-0.00051	0.00087	-0.00093	0.00233	0.00091	-0.00114
<i>IID standard-errors in parentheses</i>						
<i>Signif. Codes: ***: 0.01, **: 0.05, *: 0.1</i>						

Table B3: What is the relationship between putative identity assigned to a department and the email address type collected from the department? The first row *White X Male* describes the proportion of departments receiving an email from a White male putative identity based on the type of email type collected from that department. These values reflect the proportions of email types of emails collected (seen in Table B1). The subsequent rows show the differential for each putative identity from the baseline of White male.

Dependent Variables: Model:	Department (1)	Chief (2)	Comm. Officer (3)	Found (4)	Records (5)	Accountability (6)
<i>Putative Identities</i>						
White × Male	0.4188*** (0.0262)	0.4701*** (0.0267)	0.0484*** (0.0105)	0.0285*** (0.0073)	0.0285*** (0.0070)	0.0057 (0.0049)
<u>Differential from White Male</u>						
White × Female	-0.0304 (0.0373)	0.0488 (0.0379)	-0.0050 (0.0149)	0.0063 (0.0104)	-0.0198** (0.0100)	<0.0000 (0.0070)
Hispanic × Male	0.0156 (0.0373)	0.0285 (0.0379)	-0.0164 (0.0149)	-0.0110 (0.0104)	-0.0168* (0.0100)	0.0001 (0.0070)
Hispanic × Female	-0.0323 (0.0369)	0.0509 (0.0376)	-0.0008 (0.0148)	-0.0145 (0.0103)	-0.0089 (0.0099)	0.0055 (0.0069)
Black × Male	-0.0342 (0.0371)	0.0855** (0.0377)	-0.0199 (0.0148)	-0.0171* (0.0103)	-0.0171* (0.0099)	0.0028 (0.0070)
Black × Female	0.0122 (0.0372)	0.0098 (0.0378)	-0.0082 (0.0149)	-0.0199* (0.0103)	-0.0026 (0.0100)	0.0087 (0.0070)
<i>Fit statistics</i>						
Observations	2,095	2,095	2,095	2,095	2,095	2,095
R ²	0.00188	0.00325	0.00146	0.00471	0.00330	0.00125
Adjusted R ²	-0.00051	0.00087	-0.00093	0.00233	0.00091	-0.00114
<i>IID standard-errors in parentheses</i>						
<i>Signif. Codes: ***: 0.01, **: 0.05, *: 0.1</i>						

The results in Table B3 demonstrate that, due in large part to random assignment of treatment, there are few statistically significant differences in the proportions of email types used for each putative identity. Only one result is noteworthy, the proportion of Black male emails sent to police chiefs.³¹ About 47% of departments that received emails from White male putative identities had a *Chief* email type. In comparison, departments that received emails from Black male putative identities, were statistically significantly more likely to be a *Chief* email type (8.55%, [1.15%, 15.94%]).

³¹White female emails sent to *Records* email types is also significant at the 5% level, but given the number of observations, this is not of much concern.

Table B4 compares response rates for the putative identities across the different department email address types. The remaining five putative identity response rates are compared to mean response rate for White male putative identities. The second column of Table B4, *Baseline*, uses all of the data from the study (n = 2095), and is used as a point of reference for the other two models.

Table B4: Different types of email addresses were collected for departments. In addition to heterogeneity of response rates across putative identities, there is additional heterogeneity across email types. Column 2, *Baseline*, are the response rate differences for each putative identity compared to the response rate for White male putative identities ignoring department email address type. Column 3 and 4 have the same structure as column 2, but are subsets of the data restricted to specific department email type.

Dependent Variable: Model:	Baseline	response Department	Chief
<i>Putative Identities</i>			
White \times Male (<i>mean</i>)	0.7580	0.7619	0.7515
White \times Female	-0.0285 (0.0479)	-0.0289 (0.0626)	-0.0064 (0.0714)
Hispanic \times Male	-0.1500*** (0.0354)	-0.1378** (0.0557)	-0.1438** (0.0594)
Hispanic \times Female	-0.0928* (0.0423)	-0.1337* (0.0721)	-0.0644 (0.0539)
Black \times Male	-0.1394*** (0.0377)	-0.2203*** (0.0615)	-0.0812 (0.0471)
Black \times Female	-0.0970** (0.0397)	-0.1073** (0.0408)	-0.0927 (0.0627)
<i>Fixed-effects</i>			
week	Yes	Yes	Yes
dept_address_state	Yes	Yes	Yes
<i>Fit statistics</i>			
Observations	2,095	853	1,063
R ²	0.06214	0.10994	0.11299
Within R ²	0.01404	0.02433	0.01212
<i>Clustered (week & dept_address_state) standard-errors in parentheses</i>			
<i>Signif. Codes: ***: 0.01, **: 0.05, *: 0.1</i>			

The estimates in third and fourth columns (*Department* and *Chief*) of Table B4 reveal considerable heterogeneity in response for some of the putative identities. Police departments that had a *Department* email

address type were much less likely to respond to Black male putative identities than police departments that had a *Chief* email address type: -22.03[-35.95, -8.11] vs. -8.12 [-18.77, 2.53]. Similarly, Hispanic female putative identity response rate differentials were twice the size in magnitude for police departments that had a *Department* email address type compared to police departments that had a *Chief* email address type: -13.37 [-29.68, 2.93] vs. -6.44 [-18.63, 5.74]. The differences across email types for the other putative identities are smaller, but in all instances, except the Hispanic male putative identities, police departments that had a *Department* email address type had smaller response rate differentials compared to the White male response rate. A small portion of this differential is driven by marginally higher White male response rates for police departments that had a *Department* email address type compared to police departments that had a *Chief* email address type (76.19% vs. 75.15%).

The heterogeneity in distribution of department email address types across the different putative identities makes interpreting the results from Table B4 slightly more challenging. However, the results do suggest that the presence of biased policing behavior may vary by the chain of command within departments. There is correlation between the local population that a department serves and which email address type is collected for a department. The bigger departments are more likely to have a departmental email address publicly available than departments that serve smaller populations. However, as seen in Figure 3 there does not seem to be an obvious relationship between local population size and response rates. Additionally, it is unclear which of the two mechanisms (population or email address type) if either of them, are driving the heterogeneous response rates. An alternative explanation is that departments that make the police chief the public point of contact are different than those that do not. However, this explanation is slightly confounded by the fact that many departments that had a publicly available departmental email address also had a email address for the police chief. Perhaps the most compelling explanation is that police chiefs are very unlikely to be the officer that will be mentioned in the eventual complaint and has the most interest in ensuring that their subordinates are maintaining a high level of professionalism. In contrast, a departmental email address is likely to be maintained by officers that may worry they will be mentioned in the complaint.

B Online Appendix: Identity Construction

Six different “identities” will be used:

1. Black Female
2. Black Male
3. White Female
4. White Male
5. Hispanic Female
6. Hispanic Male

Consistent with standard practices in the correspondence study literature, identity (gender and race/ethnicity) of the email sender will be implied by name (first name and last name). Ten unique first names and six unique last names are chosen for each identity (60 unique name combinations for each identity). Using multiple names for each identity minimizes the importance of a specific name.

- First names are selected from research done by Gaddis ([Gaddis \(2017a\)](#), [Gaddis \(2017b\)](#)). The top ten most racially identifiable first names (when coupled with last names), are chosen.
- Last names are selected from the 2010 Census. Three criteria are used to select last names:
 1. Percent of persons with that name having a specific race/ethnicity (e.g., White)
 2. Percent of persons with that name having the other relevant race/ethnicity (e.g., Black or Hispanic)
 3. The rank of the name (i.e. how common the last name is in the United States)

Name Search Equation: I selected surnames for this experiment that were both (1) racially distinctive and (2) commonly found. Priority was given to racially distinctive, because of the importance of race in the design of the experiment. However, I also wanted to avoid the scenario where police departments act differently if they see an exceedingly uncommon last name. In other words, I want race, and only race, to

be communicated by the name of the identity. The three equations below reflect the priorities I used to select the names. I decided it was unnecessary to difference the Hispanic surnames with the other two groups because of how uncommon it was for Black and White people to have a surname commonly used by Hispanic people.

- For Black Names: $percent\ race_{black} - percent\ race_{white}) - .05 \times rank_{black\ name}$
- For White Names: $percent\ race_{white} - percent\ race_{black}) - .05 \times rank_{white\ name}$
- For Hispanic Names: $percent\ race_{hispanic} - .05 \times rank_{white\ name}$

The full list of names can be inferred by the following two tables (there are 360 unique name combinations). Six high-profile recognizable celebrity names were omitted: Denzel Washington, Tyra Banks, DaShawn Jackson, Seth Meyer(s), Katelyn Olson and Pedro Martinez. These names have widespread recognition and during the testing process, respondents noted that they strongly associate these names with the celebrities having the same name.

Last Names		
White	Black	Hispanic
Olson	Washington	Hernandez
Schmidt	Jefferson	Gonzalez
Meyer	Jackson	Rodriguez
Snyder	Joseph	Ramirez
Hansen	Williams	Martinez
Larson	Banks	Lopez

First Names					
White Male	White Female	Black Male	Black Female	Hispanic Male	Hispanic Female
Hunter	Katelyn	DaShawn	Tanisha	Alejandro	Mariana
Jake	Claire	Tremayne	Lakisha	Pedro	Guadalupe
Seth	Laurie	Jamal	Janae	Santiago	Isabella
Zachary	Stephanie	DaQuan	Tamika	Luis	Esmeralda
Todd	Abigail	DeAndre	Latoya	Esteban	Jimena
Matthew	Megan	Tyrone	Tyra	Pablo	Alejandra
Logan	Kristen	Keyshawn	Ebony	Rodrigo	Valeria
Ryan	Emily	Denzel	Denisha	Felipe	Lucia
Dustin	Sarah	Latrell	Taniya	Juan	Florencia
Brett	Molly	Jayvon	Heaven	Fernando	Juanita

As mentioned, the first names were selected from [Gaddis \(2017a\)](#) and [Gaddis \(2017b\)](#). In these studies, Gaddis analyzes the correlation between the average level of the mother's education for a given first name and accuracy of perceived race and ethnicity of that name. For instance, Black names associated with lower education levels for mothers are more often perceived as Black than Black names associated with mothers with higher average education levels. In my study, while creating the identities, the associated maternal education levels documented by Gaddis are recorded in my database.

C Online Appendix: Email Account Creation

To implement this study, sender email addresses had to be created for each putative identity. Ideally, each of the 360 identities would have a unique email address. During the pre-testing process, respondents suggested that

“firstname.lastname.birthyear@mail.com” was the most realistic email address template. However, due to constraints from popular email servers (e.g., Yahoo), this was not feasible. Instead, a unique account was made for each *last name* (18 accounts in total). Due to availability, I had to be creative in creation of the email address. All of the addresses include some version of the relevant last name.³² Often included is a birth year (e.g., **Banksss.1991@mail.com**). I do not expect that the implied birth year will be a salient component of the email, but I will make a cursory examination of the role that the email sender’s apparent age plays in response rates for police departments.

³²Due to the prevalence of people with the last names chosen for the study, it was often difficult to find available addresses with the specific last name. As a result, I had to make creative decisions to create a plausible and name-relevant address. For example, “h3rnandez.1973@mail.com”.

D Online Appendix: Example Email

D.1 Email Text

The body of text for the email has been developed in consultation with other economists and a legal expert. The primary criterion in creating the right text for these emails concerned plausibility—i.e., I needed to create an email that sounded like a genuine request from a real citizen. Drafts of the email were sent to colleagues and police departments not selected for the correspondence study to assess the plausibility of the email. The body of the email message template reads as follows:

Police Department Name,

*My name is **first name** and I am interested in filing a complaint against an officer in your department. I am not sure what to do, and would like to request information on how to make a complaint. Can you please send me this information?*

sign off

full name

Where *full name* includes a first and last name, and *sign off* is randomly assigned as “Thank you!” or “Sincerely,”. The decision to exclude a “Hi” or “Hello” was based on the increased likelihood of the email being filtered as spam during the preliminary testing process mentioned above. ³³

³³There is a small concern about this email being rejected as implausible. For example, a very small police department might know everyone with whom they have recently interacted and would be able to deduce, with little effort, that the email is fabricated. A small police department might also be more likely not to respond to an email because of staffing limitations. However, because assignment of treatment (see below) is balanced across departments, estimates should remain unbiased. In future research, an alternative email to departments with a more innocuous inquiry (e.g., “Do you have a lost-and-found?”) could shed light on the matter.

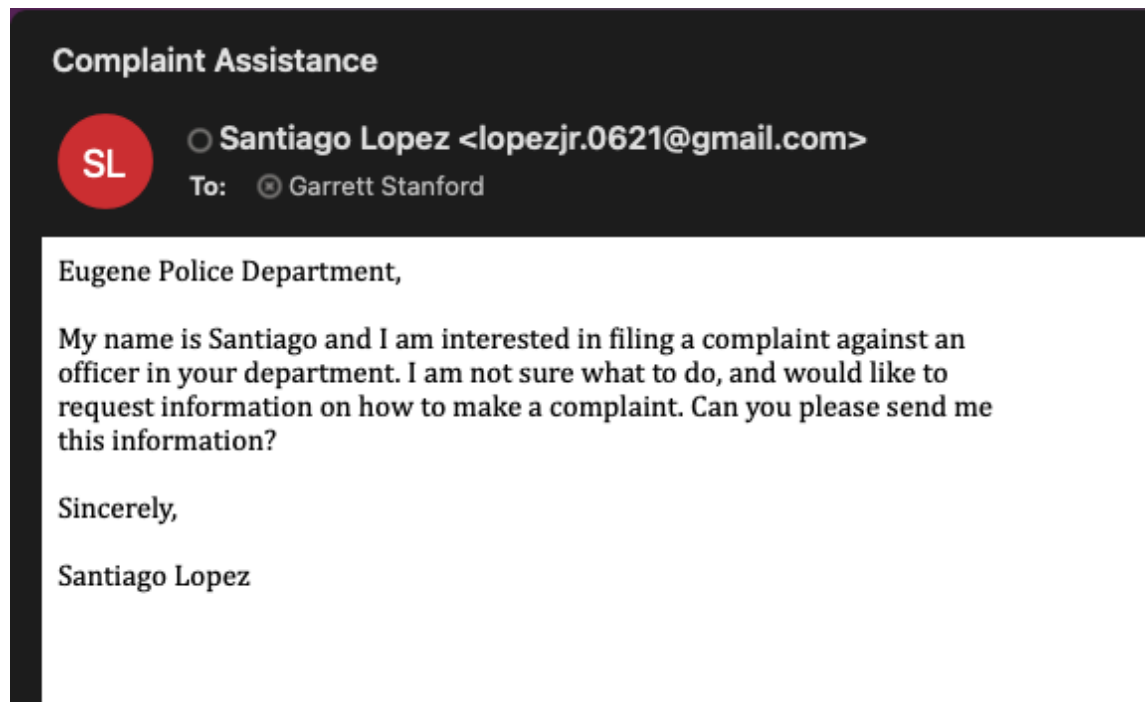


Figure B1: Example email

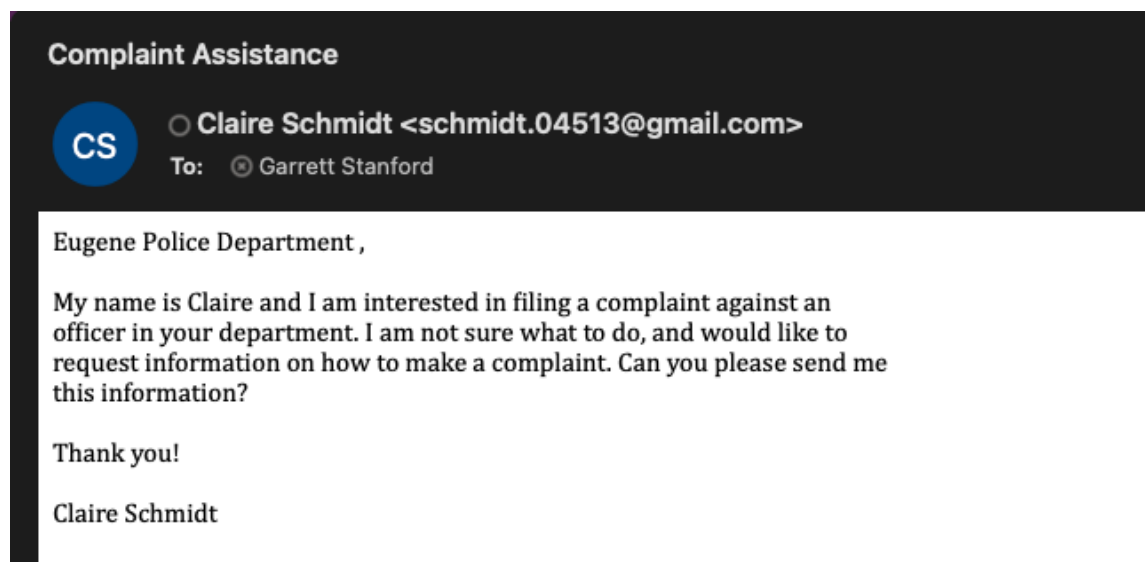


Figure B2: Example email

E Online Appendix: Treatment Assignment

Police departments are randomly assigned the sender identity they will see. The first step of treatment assignment was to balance the number of departments by state each week, so that every state received roughly the same number of emails each week. Next race and gender treatment are randomly assigned within state, with race and gender treatment levels balanced within each state. Given that assignment of emails to department by week within state was randomized, race and gender assignments are independent of week. Additionally, race and gender are roughly balanced across weeks—also as a result of the randomization of all treatment components. After week, gender and race are assigned, day of week is randomly assigned. Next, the sign off for each email is randomly assigned (the email sign off can be either “Thank you” or “Sincerely” followed by the sender’s name). The actual assignment of email sender first and last names to each department is randomized across all weeks and states.

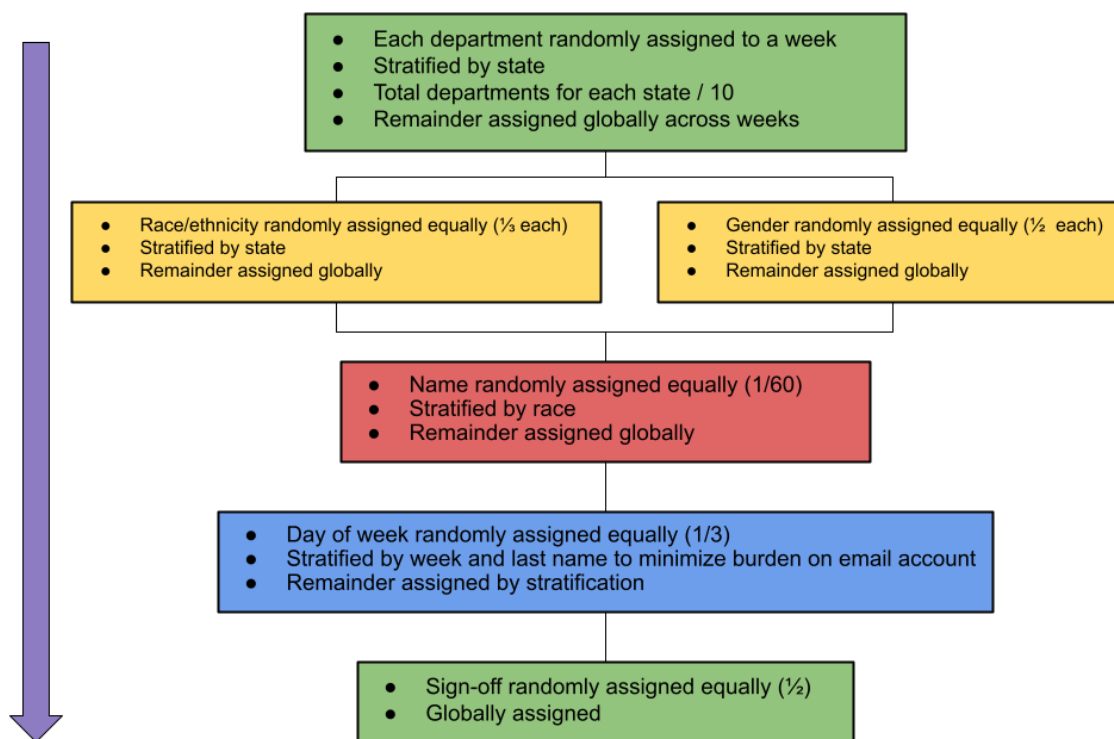


Table B1: Distribution of Race, Ethnicity and Gender identity assignment by state

State	Putative Identities					Response Statistics	
State	Black	White	Hispanic	Male	Female	Total	Mean
AK	3	2	2	4	3	5	0.714
AL	5	5	5	8	7	10	0.667
AR	6	4	6	7	9	8	0.500
AZ	6	5	4	9	6	12	0.800
CA	26	22	19	38	29	46	0.687
CO	7	7	5	10	9	14	0.737
CT	11	9	16	20	16	21	0.583
DE	2	1	1	1	3	2	0.500
FL	17	19	18	30	24	35	0.648
GA	9	10	6	10	15	14	0.560
IA	6	6	4	5	11	6	0.375
ID	3	3	5	6	5	6	0.545
IL	28	26	29	39	44	61	0.735
IN	10	10	11	16	15	19	0.613
KS	5	7	6	11	7	11	0.611
KY	3	5	4	7	5	6	0.500
LA	2	6	4	7	5	6	0.500
MA	16	17	19	28	24	31	0.596
MD	7	5	2	8	6	7	0.500
ME	6	5	6	8	9	10	0.588
MI	14	18	12	22	22	27	0.614
MN	12	8	11	13	18	24	0.774
MO	9	9	10	16	12	16	0.571
MS	6	5	5	7	9	4	0.250
MT	2	1	3	1	5	3	0.500
NC	10	6	11	12	15	17	0.630
ND	3	1	1	2	3	4	0.800
NE	4	6	2	5	7	10	0.833
NH	4	4	6	7	7	9	0.643
NJ	27	30	25	41	41	45	0.549
NM	5	3	4	7	5	3	0.250
NV	2	1	2	3	2	2	0.400
NY	14	18	19	26	25	31	0.608
OH	25	21	32	37	41	49	0.628
OK	6	7	3	8	8	10	0.625
OR	7	6	10	10	13	17	0.739
PA	25	22	21	33	35	44	0.647
RI	3	3	4	4	6	5	0.500
SC	5	6	4	10	5	7	0.467
SD	1	3	1	3	2	4	0.800
TN	6	6	6	10	8	12	0.667
TX	31	20	31	41	41	60	0.732
UT	6	4	2	8	4	9	0.750
VA	6	3	3	6	6	8	0.667

Continued on next page

Table B1 – continued from previous page

VT	2	3	3	3	5	6	0.750
WA	11	9	9	13	16	20	0.690
WI	10	10	11	17	14	26	0.839
WV	3	4	4	7	4	4	0.364
WY	1	2	3	4	2	4	0.667

Table B2: Distribution of Race, Ethnicity and Gender identity assignment by week

Week	Putative Identities					Response Statistics	
week	Black	White	Hispanic	Male	Female	Total	Mean
1	75	73	68	107	109	139	0.644
2	74	79	63	103	113	141	0.653
3	59	68	77	103	101	132	0.647
4	74	60	76	119	91	123	0.586
5	77	62	78	104	113	144	0.664
6	79	71	68	112	106	131	0.601




Figure B1: Emails sent by week.




Figure B2: Departments by state included in the study.

F Online Appendix: Experiment Implementation

I created eighteen email accounts—one for each last name. The accounts were then linked to Mozilla’s Thunderbird mail application to help automate the emailing process.³⁴ In Thunderbird, for each email address, 20 identities were created (10 females and 10 males). Although the email address that is seen by police departments cannot be arbitrarily manipulated, the “name” of the sender can be changed from message to message. For instance, an email can be sent as **Claire Olson** <olson2292@mail.com> or **Hunter Olson** <olson2292@mail.com>. This helps increase the salience of the putative identity and decrease attention to the less-specific email address itself.

Each department will receive just one email. Emails will be sent over a ten-week period. Spreading out the randomized controlled trial (RCT) over 10 weeks insures against the possibility that unique unanticipated current events could plausibly affect police department behavior (e.g., a high-profile regional or national incident involving the police). In the case of a high-profile policing incident, a weekly roll-out of the emails will allow me to detect the possible effect of any such event on police departments’ responses to the emails.

The timing of the roll-out is randomly selected using the following procedure. Police departments are randomly assigned to one of the ten weeks, while being stratified proportional to the total number of departments in each state. Each state’s total police departments (in my data set) are split into 10 equal groups and assigned to a week. In the event that, after the initial assignment, the number of departments by state are not divisible by 10, the remainder of the police departments are randomly assigned across the weeks. In the event that the total number of departments from a state is less than 10, departments are randomly assigned to the ten different weeks (with a maximum of one department per week). Each putative sender identity (i.e. email address) has the same probability of being assigned to any one of the 10 weeks.

During each week, the emails are sent on Monday, Tuesday and Wednesday. Assignment of weekday is randomized. The decision to choose different days is largely motivated by an effort to improve the ease of implementation of the emailing process for the researcher. Each email must be sent individually, so

³⁴I had originally intended to use the mailR package from R, but due to increased security policies with many popular email servers, that option is no longer as user friendly. To use mailR with, for example, Google, one needs to change the Google account settings to allow “less secure apps”. However, as of May 31st, this setting can no longer be adjusted. There are possible workarounds, but I decided to adopt an alternate strategy.

it proved easier for me to monitor the email process by spreading out the emails over a few days (with roughly 70 emails being sent each day).

All emails are sent at roughly 9 a.m. local time according to the time zone of the police department in question. However, if for a given week and given day, the same email sender address is being used for more than one police department (as dictated by the random assignment of race), a five-minute delay between each email from the same address, independent of first name, is employed. The strategy is adopted so that a single putative email account does not have to send more than one email at an exact time (i.e. at exactly 9 a.m.).

G Online Appendix: Summary Statistics

A22




Figure B1: The mean response rate by week by identities. Mean response rate across all weeks and identities (66%) is depicted by dotted black line.

A23




Figure B2: The mean response rate by week for all putative identities

H Online Appendix: Summary Statistics

Not all responses are created equally: The current analysis of the data from this correspondence study designates the outcome variable to be a police department’s timely non-automated response to a request for help. Consequently, the results are a coarse reflection of the average department’s willingness to respond to a citizen’s request for help in making a complaint about an officer. However, the premise of biased policing refers to both the frequency of interaction between officers and citizens, as well as the conduct during the interaction. Even in the specific context of an email request for a complaint form, detecting and understanding potential differences in department behavior across different sender identities is worth exploring. For example, conditional on a department providing any response, do responses differ in their helpfulness and tone across identities and, if so, how do they differ? In some instances, scrutiny of verbatim department responses reveals that not all departments are willing to guide the citizen to the officer-complaint forms. In other instances, departments specifically advise against making a formal complaint. Responses also tend to reflect a wide range of sentiment. Some departments include an apology on behalf of the department, while others simply send a phone number with no other information—the assumed implication being that the complainant should call that number for assistance. To begin to answer the question of differential response conditional response, a cursory examination of heterogeneity of responses is performed.

Table B1: Response time and word count of response measured across identities

Dependent Variables:	Word Count	Response Time (hours)
Model:	(1)	(2)
<i>Variables</i>		
White × Female	0.2948 (3.844)	-1.653 (2.573)
Hispanic × Male	3.909 (2.533)	3.032 (3.103)
Hispanic × Female	-3.051 (3.937)	-4.157 (3.056)
Black × Male	4.518 (7.728)	8.778 (7.215)
Black × Female	-4.038	2.567

Continued on next page

Table B1 – continued from previous page		
	(2.849)	(3.919)
<i>Fixed-effects</i>		
Week	Yes	Yes
State	Yes	Yes
<i>Fit statistics</i>		
Observations	1,413	1,413
R ²	0.08298	0.04623
Within R ²	0.00291	0.00892
<i>Clustered (week & dept_address_state) standard-errors in parentheses</i>		
<i>Signif. Codes: ***: 0.01, **: 0.05, *: 0.1</i>		

Table B1 reports the differentials of (1) the word count of emails from the departments and (2) the time it takes for a department respond between White male identities and the other five identities. Table B1 suggests that conditional on response, at least on the two specified dimensions, there does not seem to be any evidence of discrimination. There are a few reasons to not make strong conclusions about these null results. Most importantly, the analysis is subject to selection bias. These results are based only on the departments that *do* respond, which are different than the departments that do not respond. Additionally, word count is a crude measure of helpfulness and sentiment. An email could be helpful, friendly and to the point, but still would reflect a word count similar to an email that is unhelpful and/or unfriendly. Time of response is a stronger indicator of helpfulness. However, a quick response could be the result of a department eager to help or a department being reactive to an accusation against one of their officers. To get a strong understanding of differences in the helpfulness and sentiment of responses would require selection correction and a more rigorous sentiment analysis.